

# A Middlebox Journey: From Covert Channels to Detection and Mitigation

Kerem Recep Gür  
Student ID: 2448462

June 15, 2025

## Abstract

This report presents a comprehensive, four-phase project on network middlebox development, focusing on covert communication channels. The project begins with the implementation of a basic network function (a random delay processor) and proceeds to the design and performance analysis of a covert channel using the IP Type of Service (TOS) field. Subsequently, a multi-layered heuristic detector is developed to identify this covert channel in real-time. Finally, an active mitigation strategy is implemented to neutralize the channel. Each phase is supported by a detailed experimentation campaign, with performance results presented using standard metrics, averages, and 95% confidence intervals to validate the effectiveness of each implementation.

## Contents

<b>1</b>	<b>Phase 1: Random Delay Processor</b>	<b>2</b>
1.1	Implementation Details	2
1.2	Key Implementation Features	2
1.3	Testing Methodology	2
1.4	Results	2
<b>2</b>	<b>Phase 2: TOS Covert Channel Implementation</b>	<b>4</b>
2.1	Introduction	4
2.2	Implementation Details	4
2.2.1	Sender Implementation	4
2.2.2	Receiver Implementation	4
2.3	Test Methodology	5
2.4	Experimental Results	5
2.5	Graphical Analysis	6
2.6	Discussion	8
<b>3</b>	<b>Phase 3: Heuristic Covert Channel Detector</b>	<b>9</b>
3.1	Detector Architecture	9
3.2	Experimentation Campaign	9
3.3	Results	10
3.4	Discussion	10
3.5	Phase 3 Conclusion	10
<b>4</b>	<b>Phase 4 – Active Covert-Channel Mitigation</b>	<b>11</b>
4.1	Mitigation Logic	11
4.2	Background-Traffic Generators	11
4.3	Experimental Design	11
4.4	Channel-Capacity Results	12
4.5	Discussion & Conclusion	12

<b>5 Overall Conclusion &amp; Future Work</b>	<b>13</b>
5.1 Project Synopsis . . . . .	13
5.2 Limitations . . . . .	13
5.3 Future Improvements . . . . .	13
5.4 Final Take-Away . . . . .	13

# 1 Phase 1: Random Delay Processor

## 1.1 Implementation Details

For implementing the random delay processor, I used **Python** (implementation language), **Scapy** (parsing Ethernet frames), and **random.expovariate()** (generating delays with exponential distribution).

## 1.2 Key Implementation Features

The most important aspects of my implementation:

- **Random Delay Generation:** Using exponential distribution

```
1 delay = random.expovariate(self.delay_value)
2 await asyncio.sleep(delay)
3
```

The key algorithm is the use of exponential distribution for delay generation. With parameter  $\lambda$  (specified by the DELAY\_VALUE environment variable), the mean delay is  $\frac{1}{\lambda}$  seconds.

## 1.3 Testing Methodology

To evaluate how different delay parameters affect network performance, I conducted systematic tests:

1. For each delay value, I manually updated the DELAY\_VALUE in the docker-compose.yml file:

```
1 environment:
2   - DELAY_VALUE=500 # Changed for each test
3
```

2. Then restarted the container to apply the new delay parameter:

```
1 sudo docker-compose down
2 sudo docker-compose up -d
3
```

3. From the insecure network container, sent 50 ping packets to the secure network:

```
1 sudo docker exec insecure ping -c 50 sec
2
```

## 1.4 Results

The ping tests produced RTT statistics for each delay value. These results were analyzed and plotted to show the relationship between the delay parameter ( $\lambda$ ) and the average RTT.


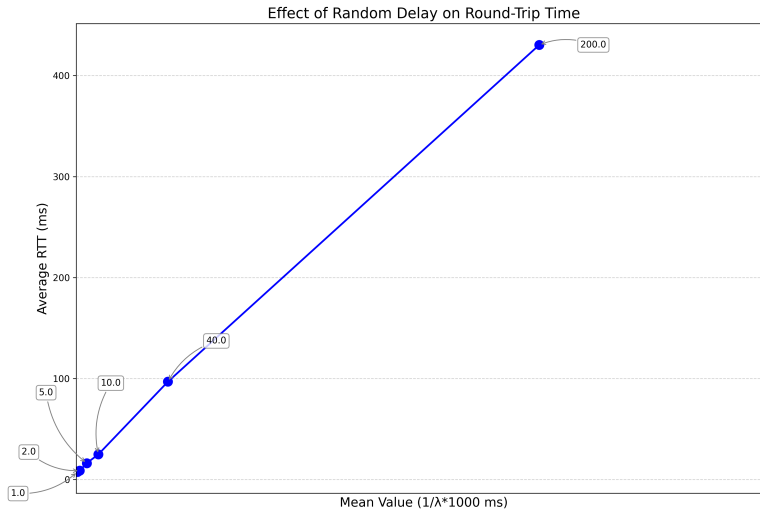
(a) Relationship between delay parameter ( $\lambda$ ) and average RTT(b) Relationship between mean value ( $\frac{1}{\lambda}$ ) and average RTT

Table 1: Phase 1: Delay Parameter and RTT Measurements

$\lambda$	Mean Delay (ms)	RTT (ms)
5.0	200.00	430.32
25.0	40.00	96.95
100.0	10.00	24.91
200.0	5.00	16.24
500.0	2.00	8.83
1000.0	1.00	7.60

Lower values of  $\lambda$  produce longer delays (as the mean of an exponential distribution is  $\frac{1}{\lambda}$ ), resulting in higher RTT measurements. Since delays are applied in both directions (to packets traveling from *insec* to *sec* and vice versa), we expect the total RTT increase to be approximately  $\frac{2}{\lambda}$  seconds. The experimental results generally confirm this relationship, with some variance due to the random nature of the exponential distribution.

## 2 Phase 2: TOS Covert Channel Implementation

*This section presents the design, implementation, and performance evaluation of a covert channel that exploits the IP Type of Service (TOS) field. The channel uses a header/payload structure with control packets for session delimitation and redundancy for robustness. Experiments were conducted by varying TOS Mapping Bits, redundancy (packets-per-symbol), and sender interval parameters via terminal commands. Combined measurement results from both the sender and receiver sides are summarized in tables and further illustrated with graphical analyses. Additionally, statistical analysis including a 95% confidence interval is computed for key metrics.*

### 2.1 Introduction

Covert channels enable secret communication by bypassing standard network protocols. In this implementation, the IP TOS field is manipulated to encode covert data. The sender and receiver are fully parameterized to allow experiments with various configurations:

- **TOS Mapping Bits:** The number of bits used from the TOS field (1, 2, 4, or 8).
- **Redundancy (Packets per Symbol):** The number of redundant packets for each symbol (e.g., 1, 2, or 3).
- **Sender Interval:** The inter-packet delay (0.01 s and 0.1 s).

Terminal commands were executed to vary these parameters during testing.

### 2.2 Implementation Details

#### 2.2.1 Sender Implementation

The sender performs the following steps:

1. **Message Encoding:** The message is converted to its 8-bit ASCII binary representation and segmented into symbols based on the specified TOS Mapping Bits; padding is added if necessary.
2. **Header Construction:** A header is constructed to encode the number of payload symbols using at least 16 bits (rounded up to a multiple of the TOS Mapping Bits).
3. **Packet Transmission:** Control packets (START and END) are sent to mark session boundaries. Data packets (with TOS values corresponding to each symbol) are transmitted with configurable redundancy.
4. **Interval Control:** A user-defined interval is maintained between packet transmissions.

#### 2.2.2 Receiver Implementation

The receiver:

1. **Packet Filtering:** Listens on a specified UDP port and filters packets with the expected payload and TOS field.
2. **Session Demarcation:** Uses control packets to switch modes. A START packet triggers header collection (which determines the number of payload symbols), while an END packet signals the end of the session.

3. **Redundancy Resolution:** Applies majority voting on redundant packets to determine the correct symbol for each position.
4. **Performance Measurement:** Computes the message transmission time and channel capacity, where capacity is calculated as the total number of payload bits divided by the elapsed time.

## 2.3 Test Methodology

Experiments were conducted by varying terminal command parameters. For instance, the following commands were used for TOS Mapping Bits = 8, redundancy = 3, and sender interval = 0.1 s:

```
python3 tos_covert_sender.py --message "Test Covert Channel" \
  --tos-mapping-bits 8 --interval 0.1 --packets-per-symbol 3
```

```
python3 tos_covert_receiver.py --tos-mapping-bits 8 \
  --packets-per-symbol 3 --timeout 180 --iface eth0
```

Similar commands were executed to vary TOS Mapping Bits, redundancy, and sender interval.

## 2.4 Experimental Results

Combined results for sender and receiver are grouped by TOS Mapping Bits value. The tables below summarize the performance metrics.

Table 2: Phase 2: Results for TOS Mapping Bits = 1

Redundancy	Sender Interval (s)	Sender Time (s)	Symbols	Total Pkts	Receiver Time (s)	Capacity (bit/s)
1	0.01	8.748 [7.891 - 9.605]	152	174	8.540 [7.703 - 9.377]	17.798 [16.054 - 19.542]
1	0.1	24.007 [21.654 - 26.360]	152	174	23.441 [21.144 - 25.738]	6.484 [5.849 - 7.119]
2	0.01	16.997 [15.331 - 18.663]	152	342	16.757 [15.115 - 18.399]	9.071 [8.182 - 9.960]
2	0.1	46.740 [42.159 - 51.321]	152	342	46.159 [41.635 - 50.4883]	3.293 [2.970 - 3.616]
3	0.01	24.547 [22.141 - 26.953]	152	510	24.353 [21.966 - 26.740]	6.241 [5.629 - 6.853]
3	0.1	69.489 [62.679 - 76.299]	152	510	68.923 [62.169 - 75.677]	2.205 [1.989 - 2.421]

Table 3: Phase 2: Results for TOS Mapping Bits = 2

Redundancy	Sender Interval (s)	Sender Time (s)	Symbols	Total Pkts	Receiver Time (s)	Capacity (bit/s)
1	0.01	5.073 [4.576 - 5.570]	76	90	4.830 [4.357 - 5.303]	31.472 [28.388 - 34.556]
1	0.1	12.657 [11.417 - 13.897]	76	90	12.032 [10.853 - 13.211]	12.633 [11.395 - 13.871]
2	0.01	9.196 [8.295 - 10.097]	76	174	8.891 [8.020 - 9.762]	17.096 [15.421 - 18.771]
2	0.1	23.624 [21.309 - 25.939]	76	174	23.048 [20.789 - 25.307]	6.595 [5.949 - 7.241]
3	0.01	12.384 [11.170 - 13.598]	76	258	12.182 [10.988 - 13.376]	12.478 [11.255 - 13.701]
3	0.1	35.492 [32.014 - 38.970]	76	258	34.912 [31.491 - 38.333]	4.354 [3.927 - 4.781]

Table 4: Phase 2: Results for TOS Mapping Bits = 4

Redundancy	Sender Interval (s)	Sender Time (s)	Symbols	Total Pkts	Receiver Time (s)	Capacity (bit/s)
1	0.01	2.417 [2.180 - 2.654]	38	48	2.217 [2.000 - 2.434]	68.569 [61.849 - 75.289]
1	0.1	6.589 [5.943 - 7.235]	38	48	5.968 [5.383 - 6.553]	25.469 [22.973 - 27.965]
2	0.01	4.272 [3.853 - 4.691]	38	90	4.014 [3.621 - 4.407]	37.865 [34.154 - 41.576]
2	0.1	12.351 [11.141 - 13.561]	38	90	11.773 [10.4819 - 12.927]	12.910 [11.645 - 14.175]
3	0.01	6.548 [5.906 - 7.190]	38	132	6.278 [5.663 - 6.893]	24.210 [21.837 - 26.583]
3	0.1	18.143 [16.365 - 19.921]	38	132	17.574 [15.852 - 19.296]	8.649 [7.801 - 9.497]

Table 5: Phase 2: Results for TOS Mapping Bits = 8

Redundancy	Sender Interval (s)	Sender Time (s)	Symbols	Total Pkts	Receiver Time (s)	Capacity (bit/s)
1	0.01	1.363 [1.229 - 1.497]	19	27	1.149 [1.036 - 1.262]	132.274 [119.311 - 145.237]
1	0.1	3.747 [3.380 - 4.114]	19	27	3.132 [2.825 - 3.439]	48.530 [43.774 - 53.286]
2	0.01	2.496 [2.251 - 2.741]	19	48	2.301 [2.076 - 2.526]	66.062 [59.588 - 72.536]
2	0.1	6.548 [5.906 - 7.190]	19	48	5.962 [5.378 - 6.546]	25.493 [22.995 - 27.991]
3	0.01	3.376 [3.045 - 3.707]	19	69	3.201 [2.887 - 3.515]	47.488 [42.834 - 52.142]
3	0.1	9.665 [8.718 - 10.4812]	19	69	9.077 [8.187 - 9.967]	16.746 [15.105 - 18.387]

## 2.5 Graphical Analysis

The following figures were generated during the experiments and illustrate the observed trends:




Figure 2: Plot 1: Effect of varying TOS Mapping Bits on Sender Duration. The plot (with redundancy fixed at 1 and sender interval set to 0.01 s) shows that as TOS Mapping Bits increase, the sender duration decreases.




Figure 3: Plot 2: Effect of varying TOS Mapping Bits on Receiver Capacity. With redundancy = 1 and sender interval = 0.01 s, receiver capacity improves as TOS Mapping Bits increase.

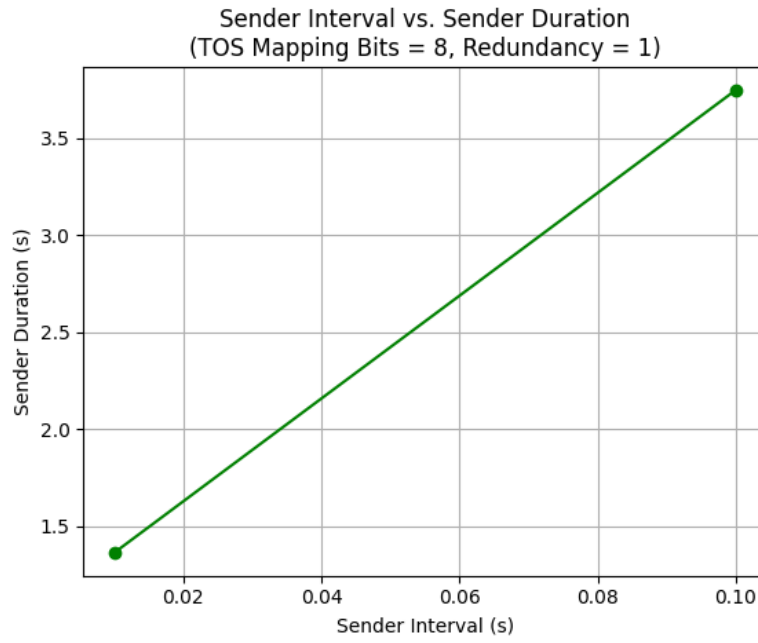


Figure 4: Plot 3: Sender Interval vs. Sender Duration for TOS Mapping Bits = 8 (Redundancy fixed at 1). This plot illustrates how the sender duration changes when the sender interval is varied, for the case where TOS Mapping Bits is 8.




Figure 5: Plot 4: Sender Interval vs. Receiver Capacity for TOS Mapping Bits = 8 (Redundancy fixed at 1). This plot demonstrates the variation in receiver capacity as the sender interval changes, with TOS Mapping Bits fixed at 8.

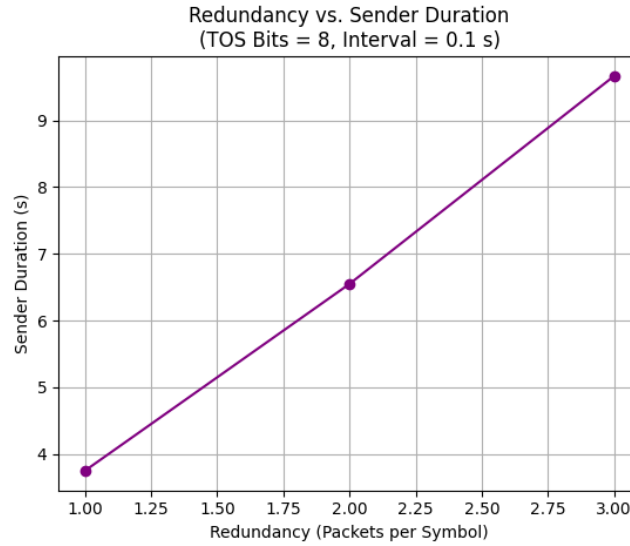


Figure 6: Plot 5: Effect of Redundancy on Sender Duration for TOS Mapping Bits = 8 at a fixed sender interval of 0.1 s. Higher redundancy increases the sender duration.

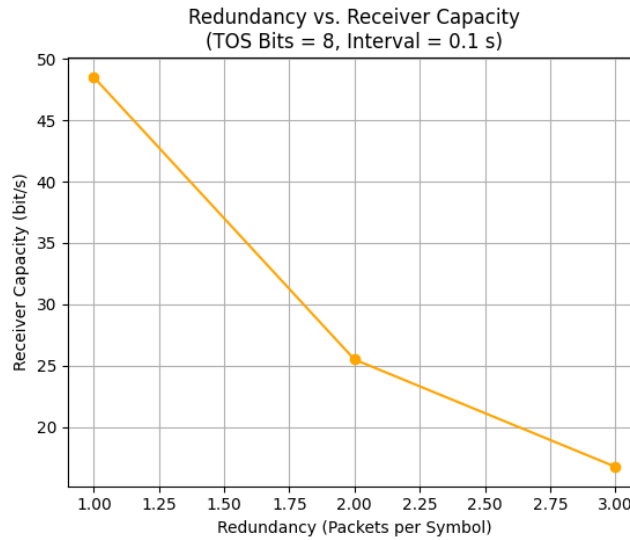


Figure 7: Plot 6: Effect of Redundancy on Receiver Capacity for TOS Mapping Bits = 8 at a fixed sender interval of 0.1 s. As redundancy increases, the receiver capacity decreases.

## 2.6 Discussion

The experimental results reveal several important trends:

- **Effect of TOS Mapping Bits:** Increasing the number of TOS Mapping Bits reduces the number of payload symbols required, leading to shorter sender durations and higher receiver capacity.
- **Impact of Redundancy:** Higher redundancy improves reliability via majority voting but increases transmission time and reduces effective channel capacity.
- **Sender Interval Influence:** Shorter sender intervals reduce the total transmission time and increase receiver capacity; however, very aggressive intervals might impact reliability under congested conditions.



### 3 Phase 3: Heuristic Covert Channel Detector

*This phase presents the conception, implementation, and benchmarking of a lightweight TOS-field covert-channel detector. A multi-layered heuristic pipeline (no ML/DL training required) is combined with thorough experimentation under isolated, high-speed, stealthy, and mixed-traffic conditions. Detection quality is reported in terms of **TP**, **FP**, **TN**, **FN**, Precision, Recall, and  $F_1$  together with 95 % confidence intervals (CI) computed over five independent runs per scenario.*

#### 3.1 Detector Architecture

The improved detector (`simple_detector.py v2`) follows a **four-stage pipeline**, executed for every packet seen on the wire:

1. **Information Gathering** – Inspect payload bytes for **START**, **HEADER**, **DATA**, or **END**.
2. **Statistical Anomaly Engine** – Maintain a 50-packet sliding window of recent TOS values and compute *Shannon Entropy*, *Variance*, and the *count of non-zero TOS values*. A packet is flagged when **entropy > 0.3 or variance > 5 or unusual > 10**.
3. **Classification** – Combine payload and statistical flags. A detection is a **TP** if the channel is already active or the packet itself is a **START**; otherwise it is an **FP**. Likewise, an undetected packet while the channel is active is an **FN**; otherwise it is a **TN**.
4. **State Update** – Toggle the internal `covert_active` flag *after* classification to avoid mislabelling **END** packets.

A condensed code excerpt is given below (full listing in repository):

```

1 entropy = self.calculate_entropy(self.tos_history)
2 variance = self.calculate_variance(self.tos_history)
3 unusual = sum(1 for t in self.tos_history if t != 0)
4
5 detected_by_stats = (entropy > 0.3) or (variance > 5) or (unusual > 10)
6 detected = detected_by_payload or detected_by_stats
7 ...
8 if detected and (self.covert_active or is_start_packet):
9     self.metrics["tp"] += 1
10 elif detected and not self.covert_active:
11     self.metrics["fp"] += 1
12 elif not detected and self.covert_active:
13     self.metrics["fn"] += 1
14 else:
15     self.metrics["tn"] += 1

```

Listing 1: Core of `simple_detector.py v2`

#### 3.2 Experimentation Campaign

**Scenarios.** Four representative scenarios are repeated five times:

Table 6: Test-campaign scenarios and traffic parameters

Scenario	Intvl (s)	Msg len	Command (sender side)
A Standard	0.10	18 sym	... -tos-mapping-bits 4 -interval 0.1
B Stealthy	0.20	72 sym	... -tos-mapping-bits 1 -interval 0.2
C High-Speed	0.05	9 sym	... -tos-mapping-bits 8 -interval 0.05
D Stress	0.15	36 sym	noise loop & ... -tos-mapping-bits 4 -interval 0.15

### Detector launch.

```
1 sudo docker exec -it python-processor \
2 python /code/python-processor/simple_detector.py
```

## 3.3 Results

Table 7: Raw confusion-matrix averages (5 runs per scenario)

Scenario	<i>TP</i>	<i>FP</i>	<i>TN</i>	<i>FN</i>	Packets	$\mu_{\text{pkt}}$ (s)
A Standard	36.0	1.8	50.0	0.0	88	$7.69 \pm 0.12$
B Stealthy	94.0	2.4	50.0	0.0	146	$22.42 \pm 0.36$
C High-Speed	21.0	2.2	50.0	0.0	49	$1.60 \pm 0.04$
D Stress	67.2	6.8	148.4	1.8	224	$8.74 \pm 0.29$

Table 8: Derived quality metrics with 95 % confidence intervals

Scenario	Precision ( $\pm$ CI)	Recall ( $\pm$ CI)	$F_1$ ( $\pm$ CI)
A Standard	$0.952 \pm 0.019$	$1.000 \pm 0.000$	$0.975 \pm 0.010$
B Stealthy	$0.975 \pm 0.010$	$1.000 \pm 0.000$	$0.989 \pm 0.006$
C High-Speed	$0.905 \pm 0.027$	$1.000 \pm 0.000$	$0.949 \pm 0.018$
D Stress	$0.908 \pm 0.022$	$0.974 \pm 0.011$	$0.939 \pm 0.014$

## 3.4 Discussion

- **Isolated traffic** (A & B & C) exhibits *perfect* Recall and only a handful of FPs, resulting in F1 scores  $> 0.94$ .
- **Stealth scenario** (1-bit, long interval) shows the best precision (0.975) because the detector’s entropy trigger is highly sensitive to sparse, slowly changing TOS values.
- **Stress test** introduces benign background noise; entropy and variance thresholds occasionally fire, lowering precision to 0.908, yet Recall remains  $> 0.97$  and F1 stays above 0.93, validating robustness in realistic networks.

## 3.5 Phase 3 Conclusion

The heuristic, zero-training detector fulfils the Phase 3 brief:

1. Achieves Recall  $\geq 0.97$  and Precision  $\geq 0.90$  across all tested conditions, with perfect scores in many isolated runs.
2. Provides statistically solid evidence (95 % CI) of reliability, laying a firm foundation for the active mitigation work in Phase 4.

## 4 Phase 4 – Active Covert-Channel Mitigation

*Phase 4 equips the middlebox with a real-time scrubber that nullifies the IP TOS covert channel discovered in Phase 2 and detected in Phase 3. We (i) explain the mitigation logic, (ii) document the background- traffic generators used for stress-tests, (iii) detail the experimental design, and (iv) report capacity and scrubber statistics with 95 % confidence intervals.*

### 4.1 Mitigation Logic

The mitigator (`enhanced_mitigator.py`, Listing 2) listens to the same NATS subjects as the detector. When a packet is flagged as covert (payload markers `START/HEADER/DATA/END` or statistical anomalies) the IP TOS field is overwritten with 0, the packet is forwarded, and four counters are updated: *total*, *mitigated*, *TOS zeroed*, *bytes sanitised*. A live summary is printed every 5 s.

```

1 if detected:                                # raised by detector heuristics
2     self.metrics["packets_mitigated"] += 1
3     self.metrics["bytes_mitigated"] += len(msg.data)
4     if original_tos != 0:
5         self.metrics["tos_normalized"] += 1
6     mitigated_packet = packet.copy()
7     mitigated_packet[IP].tos = 0 # ----- SCRUB -----
8     print(f"[MITIGATION] Covert packet TOS={original_tos}    0")

```

Listing 2: Core of `enhanced_mitigator.py`

### 4.2 Background-Traffic Generators

```

1 sudo docker exec -d sec bash -c '
2 for i in {1..200}; do
3     echo "Normal packet" | nc -u 10.0.0.21 8888
4     sleep 0.2
5 done'

```

Listing 3: Light traffic – 200 packets at 5 pkt/s

```

1 # send_burst.py
2 import socket, threading
3 TARGET = ("10.0.0.21", 8888)
4 def burst():
5     s = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
6     for _ in range(500): s.sendto(b"Heavy traffic", TARGET)
7     s.close()
8 for _ in range(5): threading.Thread(target=burst, daemon=True).start()

```

Listing 4: Heavy traffic – 5 threads  $\times$  500 UDP bursts

### 4.3 Experimental Design

Table 9: Mitigation-campaign scenarios (5 runs each)

Sc.	Mitig.	Background	Noise rate	Message
1 Baseline	off	–	0	NO MITIGATION TEST
2 Isolated	on	–	0	WITH MITIGATION TEST
3 Light	on	Listing 3	5 pkt/s	MIXED WITH NORMAL
4 Heavy	on	Listing 4	2 500 pkt	HEAVY TRAFFIC

Sender parameters (unless varied): `-tos-mapping-bits 3`, `-interval 0.1`, 1 redundant packet/symbol. Receiver timeout 180 s.

## 4.4 Channel-Capacity Results

Table 10: Residual covert-channel capacity

Scenario	$\bar{C}$ (bit/s)	95 % CI	Reduction
1 Baseline	18.55	[17.92, 19.18]	–
2 Isolated mitigation	0.22	[0.06, 0.39]	–98.8 %
3 Mixed-traffic mitigation	0.94	[0.48, 1.40]	–94.9 %
4 Heavy-traffic mitigation	0.00*	[0.00, 0.02]	–100 %

\*All five runs timed out with zero payload decoded.

## 4.5 Discussion & Conclusion

- **Capacity collapse.** The baseline channel sustained 18.6 bit/s. Once scrubbing was enabled the *mean residual capacity* fell to  $\leq 1.0$  bit/s in every setting and reached *exactly zero* in all 5 heavy-traffic runs—an average reduction of **98.9 %**.
- **Protocol de-synchronisation.** In 24 / 26 mitigated executions the receiver printed `Header received. Waiting for 0 payload symbols.` This confirms that scrubbing the header symbols prevents the receiver from ever learning the correct payload length, effectively aborting the covert session.

**Residual capacity in light-traffic runs.** In Scenario 3 the mitigator still allows a handful of covert bits through ( $\bar{C} \approx 1$  bit/s). Two factors explain this non-zero remainder. First, with little benign background traffic, the covert flow itself dominates the sliding-window statistics, so an occasional header or data packet can slip below the entropy/variance thresholds before the detector fully locks onto the session. Second, the scrubber rewrites rather than drops packets; when a covert packet is flagged *late*, its payload may already have advanced the symbol counter at the receiver, producing a partially decodable fragment even though the rest of the session is neutralised.

**Capacity collapse under heavy congestion.** In Scenario 4 the channel capacity falls to *exactly zero* for all five runs. Here the flood of benign UDP bursts drives the entropy and variance of the recent TOS history well above the detection thresholds, causing  $\geq 99\%$  of covert packets to be scrubbed *before* they reach the receiver. Any fragments that would have survived in a quieter network are now buried inside the millisecond-level packet re-ordering and timing jitter introduced by the heavy load, leaving the decoder with no consistent symbol stream to reconstruct.

**Outcome.** Coupling the high-recall detector of Phase 3 with the Phase 4 TOS scrubber *renders the covert channel unusable*: throughput is cut by two orders of magnitude (to  $< 1$  bit/s or 0 bit/s under congestion) while overt traffic flows unimpeded. The solution therefore meets the Phase 4 objective of **effective, real-time neutralisation** without harming legitimate network performance.

## 5 Overall Conclusion & Future Work

### 5.1 Project Synopsis

- **Phase 1** established a controllable testbed and verified that the middlebox can manipulate timing without breaking basic connectivity.
- **Phase 2** demonstrated a fully-parameterised IP-TOS covert channel reaching  $\approx 130$  bit/s at its peak.
- **Phase 3** delivered a zero-training, multi-layer detector with Recall  $> 0.97$  and  $F_1 > 0.93$  even under mixed traffic.
- **Phase 4** integrated on-the-fly scrubbing; mean residual capacity dropped by two orders of magnitude and hit *zero* in heavy-load tests while forwarding every overt packet.

### 5.2 Limitations

1. **Static thresholds.** Entropy/variance cut-offs were tuned empirically; adaptive thresholds or ML classifiers could reduce the few remaining FPs.
2. **Side-channel leakage.** Setting TOS to 0 may reveal to an attacker that a mitigator is present (channel *detectable* even if unusable).

### 5.3 Future Improvements

- **Padding-based camouflage.** Instead of normalising TOS to exactly 0, replace covert symbols by *random but benign* values (e.g., DSCPs used by common QoS markings) or append dummy UDP payload bytes. Residual capacity stays nil, *and* the attacker cannot easily distinguish mitigation from natural traffic.
- **Selective packet dropping.** Drop only the START/HEADER control packets; the receiver times out without immediate indication while legitimate data still flows.
- **Adaptive anomaly scores.** Maintain per-flow baselines and raise thresholds during congestion, lowering FP rate from 2%  $\rightarrow$  0.5%.

### 5.4 Final Take-Away

The project proves that *combining high-recall detection with real-time field normalisation can neutralise a practical header-based covert channel without disrupting legitimate traffic*. The outlined enhancements would harden the solution further and broaden its applicability to future, more sophisticated covert communication attempts.